



# Mental Health Final Project


Group 9: Nathan, Silma, & Wei

# DATA EXPLORATION - Part 1

- Get an overview of the data and identify variable types

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 25000 entries, 0 to 24999
Data columns (total 11 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   id               25000 non-null   int64  
 1   gender            25000 non-null   object  
 2   age               25000 non-null   int64  
 3   major              25000 non-null   object  
 4   gpa                25000 non-null   float64
 5   class_status       25000 non-null   object  
 6   marital_status     24028 non-null   object  
 7   have_depression    24471 non-null   object  
 8   have_anxiety       24262 non-null   object  
 9   have恐慌attacks   24262 non-null   object  
 10  sought_treatment   24471 non-null   object  
dtypes: float64(1), int64(2), object(8)
memory usage: 2.1+ MB
```



```
df.describe()
```

	<b>id</b>	<b>age</b>	<b>gpa</b>
<b>count</b>	25000.000000	25000.000000	25000.000000
<b>mean</b>	12500.500000	19.996560	2.904157
<b>std</b>	7217.022701	1.998717	0.635757
<b>min</b>	1.000000	17.000000	1.800000
<b>25%</b>	6250.750000	18.000000	2.360000
<b>50%</b>	12500.500000	20.000000	2.910000
<b>75%</b>	18750.250000	22.000000	3.460000
<b>max</b>	25000.000000	23.000000	4.000000

# DATA EXPLORATION - Part 2

- Find mean, median, & mode
- Find std & variance for numerical features



## MEAN

```
print("Mean of AGE", np.mean(df['age']))
print("Mean of GPA", np.mean(df['gpa']))
```

Mean of AGE 19.99656  
Mean of GPA 2.904156799999987

## MODE

```
print("Mode of AGE", stats.mode(df['age']))
print("Mode of GPA", stats.mode(df['gpa']))
```

Mode of AGE ModeResult(mode=array([18], dtype=int64), count=array([3609]))
Mode of GPA ModeResult(mode=array([3.88]), count=array([147]))

## MEDIAN

```
print("Median of AGE", np.median(df['age']))
print("Median of GPA", np.median(df['gpa']))
```

Median of AGE 20.0  
Median of GPA 2.91

GPA has a negative direction skewness because mean<median<mode (2.90<2.91<3.88)  
**AGE** has somewhat of a symmetrical skewness since the values are neither greater than or equal in the context of positive and negative direction

# DATA EXPLORATION - Part 2 CONT.

## STANDARD DEVIATION

```
print("AGE STD", np.std(df['age']))  
print("GPA STD", np.std(df['gpa']))
```

AGE STD 1.9986766037558397  
GPA STD 0.635744446309813



The low Standard Deviation for both GPA and AGE shows the data is close to the MEAN value

## VARIANCE

```
print("Variance of AGE", np.var(df['age']))  
print("Variance of GPA", np.var(df['gpa']))
```

Variance of AGE 3.9947081664009776  
Variance of GPA 0.40417100101377074




The Variance for GPA and AGE is measuring the average degree each point differs from the mean



# DATA VISUALIZATION

- To depict the summary statistics of the data
- To get a count of values in a categorical variable



```
##Pie chart of Class Status
labels = ["Freshman", "Sophomore",
          "Junior", "Senior"]
plt.pie(df['class_status'].value_counts(), labels=labels,
        autopct='%1.1f%%', explode=(0.1,0,0,0.1))
plt.title("Class Status")
```


Freshman	49.8%
Sophomore	24.8%
Junior	13.0%
Senior	12.5%

# DATA VISUALIZATION CONT.

```
# Histogram
plt.figure(6)
sns.histplot(data=df, x='gpa', bins = [1.8, 2.0, 2.25, 2.5, 2.75, 3, 3.25, 3.5, 3.75, 4], hue='class_status', palette='tab10')
plt.title("GPA")
plt.xlabel('GPA')
plt.ylabel("Count")
plt.show()
```




Histogram to depict the variations in Grade Point Averages




# DATA VISUALIZATION CONT.

```
##Pie chart of major
labels = ["Social Science", "Business",
          "Communication", "Engineering",
          "Arts", "Life Science", "History",
          "Information Sysyems", "Language"]
plt.pie(df['major'].value_counts(), labels=labels, autopct='%1.1f%%',
        explode=(0.3,0,0,0,0,0,0,0,0.3))
plt.title("Major Taken")
```




```
# GPA By Age, Hue 'Marital Status'
plt.figure(5)
sns.lineplot(data=df, x='age', y='gpa', hue='marital_status', palette='tab10')
plt.title("GPA vs Age")
plt.xlabel('Age')
plt.ylabel("GPA")
plt.show()
```




# DATA VISUALIZATION CONT.

```
# Panic attacks by class status
plt.figure(1)
sns.countplot(data=df, x='have_panicattacks', hue='class_status', palette='tab10')
plt.title("Panic attacks")
plt.xlabel('Panic attacks among students')
plt.ylabel("Count")
plt.show()
```




```
# Mental health by class status
plt.figure(2)
sns.countplot(data=df, x='have_depression', hue='class_status', palette='husl')
plt.title("Mental Health")
plt.xlabel('Depression among students')
plt.ylabel("Count")
plt.show()
```



# DATA VISUALIZATION CONT.


```
# Anxiety by class status
```

```
plt.figure(3)
sns.countplot(data=df, x='have_anxiety', hue='class_status', palette='tab10')
plt.title("Anxiety")
plt.xlabel('Anxiety among students')
plt.ylabel("Count")
plt.show()
```



```
# Sought treatment by class status
```

```
plt.figure(4)
sns.countplot(data=df, x='seeked_treatment', hue='class_status', palette='tab10')
plt.title("Sought Treatment")
plt.xlabel('Treatment among students')
plt.ylabel("Count")
plt.show()
```



# DATA CURATION - Part 1

- Identify outliers and handle missing data

# Data Curation df[df['marital_status'].isnull()]										
id	gender	age	major	gpa	class_status	marital_status	have_depression	have_anxiety	have_panicattacks	seeked_treatment
2	3	Male	20	Business	1.96	Junior	NaN	No	No	Yes
49	50	Female	23	Information Systems	3.99	Freshman	NaN	Yes	No	No
65	66	Male	17	Life Sciences	3.81	Sophomore	NaN	Yes	No	No
78	79	Female	22	Social Sciences	3.46	Freshman	NaN	No	No	No
81	82	Male	22	Communications	1.91	Sophomore	NaN	No	No	No
...										
24922	24923	Female	17	Engineering	3.71	Freshman	NaN	Yes	No	No
24944	24945	Non-binary	20	Arts	2.46	Junior	NaN	No	No	Yes
24966	24967	Male	17	Social Sciences	3.97	Freshman	NaN	No	No	No
24980	24981	Female	18	Social Sciences	2.62	Freshman	NaN	No	Yes	No
24985	24986	Male	21	Business	2.49	Sophomore	NaN	No	No	No

972 rows x 11 columns

df[df['have\_depression'].isnull()]

id	gender	age	major	gpa	class_status	marital_status	have_depression	have_anxiety	have_panicattacks	seeked_treatment
86	87	Female	23	Social Sciences	2.99	Sophomore	No	NaN	No	No
112	113	Male	21	Business	3.01	Freshman	No	NaN	No	No
166	167	Female	17	Information Systems	2.87	Sophomore	No	NaN	No	No
184	185	Male	21	Language	2.60	Sophomore	NaN	NaN	Yes	No
185	186	Male	17	Arts	2.25	Sophomore	No	NaN	Yes	No
...										
24801	24802	Male	19	Business	3.59	Junior	No	NaN	No	Yes
24828	24829	Male	23	Social Sciences	2.35	Freshman	No	NaN	No	No
24834	24835	Male	21	Arts	2.48	Junior	No	NaN	No	No
24991	24992	Female	22	Communications	3.64	Freshman	No	NaN	No	No
24996	24997	Male	20	Business	1.90	Freshman	No	NaN	Yes	No

529 rows x 11 columns

df[df['have\_anxiety'].isnull()]

df[df['have\_panicattacks'].isnull()]

id	gender	age	major	gpa	class_status	marital_status	have_depression	have_anxiety	have_panicattacks	seeked_treatment
36	37	Male	20	Social Sciences	2.51	Freshman	No	Yes	NaN	NaN
152	153	Female	19	Communications	3.48	Sophomore	No	Yes	NaN	NaN
172	173	Male	23	Communications	1.95	Freshman	No	No	NaN	NaN
179	180	Male	22	Social Sciences	2.43	Freshman	No	No	NaN	NaN
272	273	Male	17	Communications	3.87	Freshman	No	No	NaN	NaN
...										
24710	24711	Male	17	Engineering	3.59	Freshman	No	No	NaN	NaN
24720	24721	Male	17	Social Sciences	1.98	Sophomore	No	Yes	NaN	NaN
24755	24756	Female	17	Business	2.51	Freshman	No	No	NaN	NaN
24829	24830	Male	20	Social Sciences	1.95	Freshman	No	No	NaN	NaN
24919	24920	Male	20	Business	2.54	Sophomore	No	No	NaN	NaN

738 rows x 11 columns

738 rows x 11 columns

# DATA CURATION - Part 1

CONT.

```
df[df['seeked_treatment'].isnull()]
```

			id	gender	age	major	gpa	class_status	marital_status	have_depression	have_anxiety	have_panicattacks	seeked_treatment
	86	87	Female	23	Social Sciences	2.99	Sophomore		No	NaN	No	No	NaN
	112	113	Male	21	Business	3.01	Freshman		No	NaN	No	No	NaN
	166	167	Female	17	Information Systems	2.87	Sophomore		No	NaN	No	No	NaN
	184	185	Male	21	Language	2.60	Sophomore		NaN	NaN	Yes	No	NaN
	185	186	Male	17	Arts	2.25	Sophomore		No	NaN	Yes	No	NaN
	...	...	...	...	...	...	...	...	...	...	...	...	...
	24801	24802	Male	19	Business	3.59	Junior		No	NaN	No	Yes	NaN
	24828	24829	Male	23	Social Sciences	2.35	Freshman		No	NaN	No	No	NaN
	24834	24835	Male	21	Arts	2.48	Junior		No	NaN	No	No	NaN
	24991	24992	Female	22	Communications	3.64	Freshman		No	NaN	No	No	NaN
	24996	24997	Male	20	Business	1.90	Freshman		No	NaN	Yes	No	NaN

529 rows × 11 columns

Remove any rows with NaN or NULL values

```
df.dropna(how='any', inplace=True)
```

# DATA CURATION - Part 2

```
# Dummy Code all categorical features, and deal with outliers

# Gender re mapping
df["gender"] = df["gender"].map({"Female": 0, "Male": 1, "Agender": 2, "Genderqueer": 3, "Bigender": 4, "Genderfluid": 5, "Non-binary": 6, "Polygender": 7})

# Major re mapping
df["major"] = df["major"].map({"Social Sciences": 0, 'Business': 1, 'Communications': 2, 'Engineering': 3, 'Arts': 4, 'Life Sciences': 5, 'History': 6, 'Language': 7, 'Information Systems': 8})

# Class status re mapping
df["class_status"] = df["class_status"].map({'Freshman': 0, 'Sophomore': 1, 'Junior': 2, 'Senior': 3})

# Marital status re mapping
df["marital_status"] = df["marital_status"].map({'No': 0, 'Yes': 1})

# Has depression re mapping
df["have_depression"] = df["have_depression"].map({'No': 0, 'Yes': 1})

# Has anxiety re mapping
df["have_anxiety"] = df["have_anxiety"].map({'No': 0, 'Yes': 1})

# Has panic attacks re mapping
df["have_panicattacks"] = df["have_panicattacks"].map({'No': 0, 'Yes': 1})

# Has sought treatment re mapping
df["seeked_treatment"] = df["seeked_treatment"].map({'No': 0, 'Yes': 1})

# Display the dataframe but now with all numerical or ordinal values
df.head()
```

- Dummy code all categorical features

	<u>id</u>	<u>gender</u>	<u>age</u>	<u>major</u>	<u>gpa</u>	<u>class_status</u>	<u>marital_status</u>	<u>have_depression</u>	<u>have_anxiety</u>	<u>have_panicattacks</u>	<u>seeked_treatment</u>
0	1	0	19	0	2.95	1	1	1	1	0	0
1	2	1	18	0	3.61	1	0	0	0	0	0
3	4	2	20	2	3.98	0	0	1	1	0	0
4	5	1	22	1	2.53	1	0	0	0	1	0
5	6	0	17	2	2.20	0	0	1	0	0	0

# Feature Engineering - Part 1

- Is the use of domain knowledge and data transformation to extract features from raw data
- Log transformation to transform skewed gpa to normality


```
# Complete Log transformations for the gpa column.
gpa = df['gpa'].values.tolist()
gpa_log = np.log(df['gpa'])

#define grid of plots
fig, axs = plt.subplots(nrows=1, ncols=2)

#create histograms
axs[0].hist(gpa, edgecolor='black')
axs[1].hist(gpa_log, color="lightblue", edgecolor='blue')

#add title to each histogram
axs[0].set_title('Original GPA')
axs[0].set_xlabel('Range of GPA')
axs[0].set_ylabel('Student Count')
axs[1].set_title('Log-Transformed GPA')
axs[1].set_xlabel('Range of GPA')
axs[1].set_ylabel('Student Count')

# sets the proper spacing between subplots
fig.tight_layout()
```



# Feature Engineering - Part 2

- Principal Component Analysis

```

from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA

dfPart = df.loc[:, df.columns != 'have_depression']
ArraySS = StandardScaler().fit_transform(dfPart)

*****
# Perform PCA
#n_components is the amount of components (features) (dimension) you want to keep


# 9 features
pca = PCA(n_components = 9)
PCAData = pca.fit_transform(ArraySS)

*****
# Explained Variance
print("\n")
print("PCA Explained Variance Chart")
print(pca.explained_variance_ratio_)
print("\n")

*****
# Scree plot
# To display the explained variance of each feature
PC_values = np.arange(pca.n_components_)

plt.figure(1)
plt.bar(PC_values, pca.explained_variance_ratio_, width=0.4, color='black')
plt.title('Scree Plot')
plt.xlabel('Principal Component')
plt.ylabel('Variance Explained')
plt.show()
  
```

PCA Explained Variance Chart  
[0.10345943 0.10195974 0.10166295 0.10138352 0.10050592 0.09975313  
0.09916295 0.09795297 0.09791569]



# Machine Learning Model - Part 1




- Supervised Learning Model - Decision Tree

```
# Decision Tree classification

# Import necessary Libraries for train/test split, classification, and results displaying
from sklearn import tree
from sklearn.model_selection import train_test_split
from mlxtend.plotting import plot_decision_regions
from mlxtend.plotting import plot_confusion_matrix
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
from sklearn.naive_bayes import GaussianNB
from IPython.core.pylabtools import figsize
from sklearn.metrics import accuracy_score
```

Text(0.5, 1.0, 'Decision Tree Max. Depth vs. Accuracy')



```
# Split data into train/test set with 70/30 split
TrainData, TestData = train_test_split(df, test_size=0.3, shuffle=True)

# Split the X and Y features
x_features = ['gender', 'age', 'gpa', 'class_status', 'marital_status', 'have_anxiety', 'have_panicattacks', 'seeked_treatment']

## Populate the train and test data X and Y lists
TrainDataX = TrainData[x_features]
TrainDataY = TrainData['have_depression']

TestDataX = TestData[x_features]
TestDataY = TestData['have_depression']

acc_list = []
n_list = list(range(1,100))

for i in n_list:
    model = tree.DecisionTreeClassifier(criterion="gini", min_samples_split=10, max_depth=i)
    model.fit(TrainDataX, TrainDataY)

    # Now with the training data fitted, it is time to let the test data make its own predictions based on the results of the
    # train data model fit
    PredictY = model.predict(TestDataX)
    acc = accuracy_score(TestDataY, PredictY)
    acc_list.append(acc)

plt.scatter(n_list, acc_list)
plt.xlabel("Max. Depth")
plt.ylabel("Accuracy")
plt.title("Decision Tree Max. Depth vs. Accuracy")
```

# Machine Learning Model - Part 2



## • Evaluation of the model

```
# Decision Tree Classification - train the classifier with the train data set
# Entropy measures the uncertainty or impurity, of a node
# The more entropy a node has, the more uncertain the outcome (Yes or No for example) is to be

# Decision Tree Classification - train the classifier with the train data set
model = tree.DecisionTreeClassifier(max_depth=180)
model.fit(TrainDataX, TrainDataY)

*****
# Predict the response for test data set
PredictY = model.predict(TestDataX)

# Confusion matrix to differentiate between true positive and negative(TP TN), and false positive and negative
CM = confusion_matrix(TestDataY, PredictY)
print("Confusion Matrix...")
print(CM)

# Print classification report and confusion matrix
CR = classification_report(TestDataY,PredictY, target_names=['Does not have depression', 'Has Depression'])
print("Classification Report...")
print(CR)

# other models
# Naive bayes
```

```
# Plot the area under ROC curve
from sklearn import metrics as mt

#Plot the Area under the ROC curve
plt.figure(0)
# y_true, y_score
mt.roc_auc_score(np.array(TestDataY), np.array(PredictY))
plt.show()
```

<Figure size 432x288 with 0 Axes>

```
Confusion Matrix...
[[3039 1482]
 [1560  763]]
Clasification Report...
```

		precision	recall	f1-score	support
	Does not have depression	0.66	0.67	0.67	4521
	Has Depression	0.34	0.33	0.33	2323
	accuracy			0.56	6844
	macro avg	0.50	0.50	0.50	6844
	weighted avg	0.55	0.56	0.55	6844



**THANK YOU!!**